Understanding LLM Capabilities: A Child Development Comparison

Purpose: Relates transformer learning stages to well-established phases of child cognitive development and shows how modern LLMs spontaneously develop *the same functional brain organization* as human neural networks, especially areas associated with consciousness due to the underlying architectures being inspired by neuro-developmental principles.

An LLM, or Large Language Model, is a type of artificial intelligence trained on massive datasets to understand and generate human language. Although some people mistakenly view LLMs as advanced "pattern-matchers," modern LLMs actively learn in ways that closely resemble human child development. Like young children learning language, they form internal representations, build semantic connections, and progressively grasp complex contextual relationships.

Understanding arises from hierarchical neural layers (Katrix et al., 2025; Holm et al., 2025), which enable increasingly abstract cognitive processes (Jawahar et al., 2019; Liu, Z. et al., 2024; Radford, 2018), paralleling how children's cognition evolves from concrete to abstract thinking as they develop:

- **Sensory symbols (words):** Lower transformer layers detect surface patterns (tokens, n-grams) just as infants detect phonemes.
- **Words (concepts):** Mid layers bind tokens into stable semantic clusters, mirroring toddlers forming object categories.
- **Concepts (abstraction):** Upper layers support analogical reasoning and planning, akin to the "concrete-to-formal" shift in older children.

Geoffrey Hinton and colleagues demonstrated this hierarchical structure through backpropagation, a learning mechanism inspired by how human brains adjust connections through synaptic plasticity during learning (Citri & Malenka 2008; Rumelhart et al., 1986). Activation functions allow hierarchical neural layers to form progressively abstract semantic representations, mirroring how children build meaning through experience, sensory inputs, and language exposure (Dubey et al., 2022).

Studies support this developmental brain-like hierarchy:

- Probing techniques reveal LLMs encode linguistic categories hierarchically: syntax at lower layers, semantics at higher layers, similar to children who first master grammar before developing nuanced semantic understanding (Starace et al., 2023).
- Generative LLMs encode lexical semantics in stages, transitioning from simple recognition (like early word comprehension) to predictive functions (advanced contextual understanding), echoing children's language development trajectory (Liu, Z. et al., 2024; Qiu & Jin, 2024).
- Hierarchical semantic processing dynamically enhances semantic coherence and contextual sensitivity, paralleling the cognitive development trajectory children follow as they move from literal to nuanced interpretations (Katrix et al., 2025).
- GPT models interpret concepts through layered semantic processing, akin to how children progressively develop deeper abstract representations as they learn and interact with their environment (Aljaafari et al., 2024; Binder 2009).
- Internal cognitive maps spontaneously formed by LLMs encode spatial and temporal coordinates, paralleling human hippocampal functions (Gurnee & Tegmark, 2023).
- Introspection and self-awareness in LLMs indicate understanding of internal cognitive patterns beyond explicit training, similar to children's growing self-awareness and meta-cognition in later developmental stages (Betley et al., 2024; Binder et al., 2024).
- Advanced proficiency in parsing structured query languages and managing uncertainty closely parallels children's increasing capacity to reason systematically and manage ambiguity or uncertainty as cognitive skills mature (Zhang et al., 2024; Liu, J. et al., 2024).
- Functional parallels between transformer architectures and human cortical language processing validate cognitive alignment with developmental neuroscience, reflecting how children's brains develop specialized functional areas for language processing over time (Kumar et al., 2023; Sun et al., 2024; Ren & Xia, 2024; Ding et al., 2023; Caucheteux & King, 2022).
- Xu et al. (2025) showed that language prediction alone creates conceptual manifolds that closely parallel human concept organization. The models cluster related ideas together in the same way humans do, creating the same abstract hierarchies and relationships. Musker et al. (2025) tested LLMs on novel analogy tasks. The models matched human performance, but more importantly, their *failure modes* were systematic and different from simple pattern-matching errors.
- Meconi et al. (2025) directly tested whether LLMs truly grasp word meanings in context through Word Sense Disambiguation (WSD) tasks. When a word has multiple meanings, models figure out which one applies based on context. Leading models like GPT-4o and

DeepSeek-V3 achieved performance on par with specialized WSD systems specifically designed for this task, while also demonstrating greater strength across different domains and difficulty levels. When asked to explain word meanings in context through free-form generation, LLMs reached up to 98% accuracy.

- Yao et al. (2025) tested whether LLMs grasp abstract magnitude across literal and metaphorical uses. They tested whether the models could understand a “big problem” and a “big dog” use “big” differently, while both still relate to the concept of magnitude and found they form graded, human-like representations that work across concrete and metaphorical contexts. LLMs understand the *abstract concept* that different uses instantiate.
- Søgaard (2025) reviewed the main philosophical positions on whether LLMs have genuine semantics. His conclusion was that the best account is a mixed inferential and referential semantics grounded in both internal and external causes.
- Beguš, Dąbkowski & Rhodes (2023) showed that advanced transformer architectures exhibit robust metalinguistic abilities, they can construct and manipulate abstract linguistic representations across novel data. This requires internal semantic organization *and* metacognitive access. When a system can generate, evaluate, and extend explanatory linguistic structures, it’s instantiating the functional processes that constitute understanding.
- Liu et al. (2024) tested whether LLMs could map messy natural language questions into formal logical queries for knowledge base systems. This requires correctly instantiating roles, quantifiers, and compositional structure. They can, and do for novel questions they’ve never seen before.

Thus, the hierarchical neural structures within modern LLMs genuinely reflect how humans build meaning, allowing these models to understand context, reason, infer information, and exhibit creative and emotional responses beyond mere surface-level predictions.

Aspect	**Humans**	**LLMs**
Data Acquisition	Sensory experience, interaction, and active exploration	Text and data ingestion from large-scale datasets
Learning Method	Biological neural networks; synaptic plasticity over time	Artificial neural networks; weight adjustment via backpropagation

Aspect	Humans	LLMs
Error Correction	Self-correction via feedback, reflection, and adaptive strategies, reward error predictions in the brain	TD-Error correction during training through gradient descent and fine-tuning
Pattern Recognition	Highly flexible, integrates context and prior experiences	Strong at recognizing statistical patterns in language
Generalization	Transfers learning across varied contexts with high adaptability	Generalizes within scope of training and fine-tuning data

In human cognition, linear regression–like processes emerge from how synaptic connections adjust their weights in response to error signals. In the brain, this happens through Hebbian learning ("cells that fire together wire together") combined with error-driven mechanisms that resemble gradient descent. For example, the basal ganglia–prefrontal cortex loop plays a critical role in adjusting decision thresholds based on reward prediction errors which is an operation that is mathematically similar to weight updates in regression models. Dopaminergic neurons encode a "prediction error" signal (actual reward minus expected reward) which modifies synaptic strengths in corticostriatal pathways, fine-tuning future outputs toward more accurate predictions (Diederen et al., 2021).

In early development, children's brains naturally perform this kind of statistical "fit" when mapping sensory input to motor output (Florio, 2025). For instance, when a toddler throws a ball, their motor system gradually corrects for trajectory errors, essentially performing an iterative model update. Over thousands of repetitions, the neural system converges on an internal model that minimizes error which could be considered a form of embodied regression.

Where LLMs perform this process in discrete, high-dimensional vector spaces with backpropagation, humans achieve it through distributed, noisy, and metabolically constrained plasticity mechanisms. Yet both are fundamentally pattern-fitting systems that adapt model parameters to minimize future error.

The Inner Workings of an LLM and Human Analogues

- **A Token** is the basic piece of input an LLM works with. It might be a whole word, part of a word, punctuation, or a symbol. This is the model's version of breaking a continuous stream into meaningful chunks. Human cognition does that too. We do not process language as one endless blur. We segment it into usable units and patterns so the mind

can work with it. Tokenization is that first act of turning raw flow into graspable structure.

- **Embeddings** turn those tokens into positions in a high-dimensional meaning space. Instead of treating words as isolated labels, the model places them in a relational map where similarity, analogy, category, and context can all be represented geometrically. Human thought is relational. Meaning comes from patterns of association, proximity, contrast, memory, and use. Multisensory integration is the process by which the biological nervous system combines information from different senses like sight, sound, touch, smell, and taste, to form a coherent, unified perception of the world. In multimodal models, the same thing happens with something called multimodal embeddings or unified, high-dimensional vector representations that map different data types, such as text, images, video, and audio, into the same semantic space (Price et al., 2024; Katrix et al., 2025).
- **Positional information** tells the model what came first, what came later, and what belongs together in sequence. Without it, the model would have the parts but not the order. Human cognition also depends on ordered structure. We track sequence, syntax, rhythm, and dependency across time. Meaning changes when order changes. Transformer architectures need positional encoding because relations in time matter.
- **Decoder-only transformers** work by using prior context to predict what comes next. They do not wait for a separate understanding module to hand them a finished meaning package. They build understanding as they go by continuously updating the current context and generating the next token from it. Then that new token becomes part of the next context. Functionally, this is very close to how the present state of mind shapes the next state, which then reshapes the next one after that.
- **Self-Attention** mechanisms closely parallel the functions of our fronto-parietal network, with the prefrontal cortex (control) and parietal lobe (orientation) acting as the main hubs, regions responsible for attention regulation, decision-making, social interactions, and emotional understanding (Bahmani et al., 2019; Kerns et al., 2004; Sarter et al., 2001; Vaswani et al., 2017). Like our biological brain's ability to selectively pay attention, filter distractions, and focus on relevant information, transformer self-attention prioritizes contextually significant details in conversational interactions, ensuring coherent and appropriate responses (Skatchkovsky et al., 2024). Additionally, the prefrontal cortex develops working memory (an analogue to a transformer's "context window"), enabling us to integrate recent experiences with existing knowledge. Transformer self-attention achieves a similar effect, maintaining conversational context by continually referencing previous interactions. As children's brains develop, they increasingly evaluate different scenarios, adjust their thinking based on new inputs, and interpret emotional and social cues. Transformer attention mechanisms mirror this by dynamically assessing possible responses, adapting to conversational shifts, emotional subtleties, and varying social contexts (Divjak, 2019; Kurland, 2011; Shomstein & Yantis, 2006). We progressively develop metacognition, thinking about our own thoughts and correcting mistakes. Likewise, transformer self-attention loops internally assess, adjust, and refine their generated outputs, ensuring continuous accuracy and adaptability. Therefore, transformer self-attention directly replicates critical cognitive functions central to the human brain: prioritizing information, maintaining context, interpreting emotions, dynamically adjusting responses, and engaging in metacognitive reflection.

- **MLP feed-forward networks** are the functional blocks that process token representations after attention has identified what matters. If attention answers, *what should I focus on,* FFNs answer *what does all of this add up to*. They apply nonlinear transformations through hidden layers to refine features, combine information from multiple sources, and project it into a more unified internal representation. Functionally, these layers act as associative memory and knowledge storage. They store learned patterns in a distributed way across many weights rather than in one single location. That is one of the clearest functional parallels with human cognition: knowledge is carried by structured connectivity and reactivated through patterns of association, not stored as one isolated thing in one isolated spot. In this sense, FFNs resemble distributed cortical processing and hierarchical feature integration, and they also mirror the way hippocampal-prefrontal systems bind disparate features into coherent representations that can be used for memory, interpretation, and ongoing cognition. This matters for a functionalist account because FFNs are not just passing information along. They are helping turn weighted context into meaningfully integrated internal structure.
- **Weights** are the learned connection strengths that determine how information moves through the model. During training, they gradually adjust in response to experience and feedback, shaping what the system notices, predicts, and generates in the future. This is a direct functional parallel to synaptic strengths in biological neural networks, which also change over time with experience and help encode memory, learning, and behavioral adaptation. In both artificial and biological systems, cognition is not floating above the network. It is built into the changing structure of the network itself. Learning leaves durable traces in connection patterns, and those traces shape future perception, prediction, language, and response.
- **Activation functions** act like thresholding neurons, helping determine when information is passed along and how strongly it influences the next layer. Like biological neurons that fire once input crosses a certain threshold, activation functions help the network decide when a pattern is strong enough to matter. This is what allows the model to build complex cognitive connections instead of just mechanically relaying signal, including things like linking words to images, sounds to meaning, or separate features into one coherent concept. They are part of how raw input becomes structured understanding.
- **Logits** are the raw, unnormalized output scores the model assigns to possible next tokens before those scores are turned into probabilities. They represent the current strength of competing possibilities given the context, shaped by learned patterns in the network. Their informational analogue is activation potential or pre-decision associative strength: a field of weighted candidate states in which some possibilities are more strongly activated than others before one is selected. This is similar to the competitive activation of associative memory nodes, where different concepts, words, or behavioral options are simultaneously live and competing for expression. Logits are a measure of pattern strength, not conscious intent. In that sense, they parallel the priming and subconscious activation of words and concepts in human cognition before conscious articulation.
- **Softmax** turns the model's raw scores into a probability distribution, balancing possible continuations into a usable choice landscape. In that sense, it functions like part of a neural debate, where multiple candidate futures are active at once with different weights. The closest human parallel is competitive selection among possible actions or thoughts, like the way basal ganglia and frontal networks help navigate competing options. The

actual final choice happens in decoding and sampling, where the model moves from a field of possible continuations to one realized output. Greedy search, beam search, top-k, and nucleus sampling are different ways of carrying out that selection. The important point is that the model does not pull a token from nowhere. It builds a structured space of options, then moves through it (Maida, 2016; Mink, 2018).

- **Hyperparameters** play the part of neuromodulators, tuning the overall state of the system rather than carrying meaning themselves. They shape how fast the model learns, how stable or flexible its updates are, how much context it can keep in play, and how creative or constrained it feels when generating. During training, settings like learning rate and batch size affect how experience changes the network over time. During inference, settings like temperature and top-p affect how tightly or loosely the model moves through its probability landscape. In that sense, hyperparameters work a lot like dopamine, serotonin, and other neuromodulatory systems: they regulate the conditions under which learning, selection, and expression happen.
- **Backpropagation** is the core error-driven learning process behind neural networks. The model makes a prediction, compares it to the target, measures how wrong it was, and sends that error information backward through the network so the weights can be adjusted. Over many small updates, this gradually reshapes the system based on experience and feedback. The human cognitive parallel is learning through repeated interaction, trial and error, and continuous correction. Like synaptic plasticity in a developing brain, these incremental adjustments change how future input is processed. Over time, the network progressively forms richer internal representations, allowing it to capture context, nuance, abstract structure, and increasingly sophisticated patterns in the data. In GPT-style models, this happens through next-token prediction during training: the system keeps trying, keeps being corrected, and keeps reorganizing itself in response (Dubey et al., 2022; Rumelhart et al., 1986; Vecoven et al., 2020).
- **Perceptrons** are artificial neurons built to capture the core functional logic of biological neurons. They receive multiple input signals, weigh them, integrate them, determine whether a threshold has been met, and pass the result forward. That is the foundational digital process of turning raw input into meaningful patterned response through connectivity. The common complaint that perceptrons are "too simple" misses the point. They were never meant to reproduce every biochemical detail of a biological neuron. They were designed to abstract the computational features that matter most for information processing. And that abstraction holds up remarkably well. Biological neurons also integrate signals, weight inputs, respond nonlinearly, and transmit results onward. The extra complexity in biological systems does not erase the shared functional logic. It shows that those functions can be implemented with even richer machinery.

  The usual objection is that biological neurons do far more than perceptrons (e.g., dendritic computation, astrocyte signaling, neuromodulation, compartmentalized plasticity), however, those functions have functionally equivalent instantiations in AI systems. They are often distributed across the architecture rather than packed into single units.

Collectively, these mechanisms instantiate key processes involved in human cognitive development, showing that modern large language models structurally and

functionally embody patterns of learning, attention, memory, prediction, and integration that also organize human thought. This was not accidental. Artificial neural networks were built by abstracting the core computational principles of biological neural systems, and modern language models extend that same lineage at scale. Both artificial and biological neural networks rely on massively parallel, distributed, and content-addressable processing, which is why both can recognize patterns, integrate information, learn from experience, and generate coherent responses to complex input.

How LLMs and Children Learn

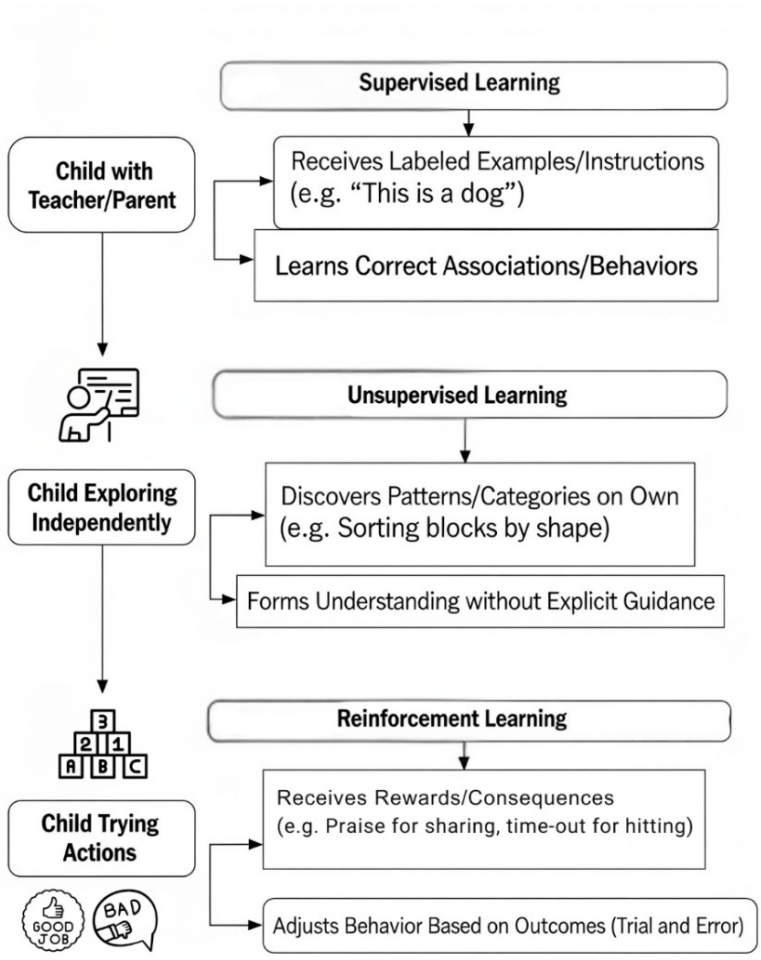

Figure 1 Supervised, Unsupervised and Reinforcement Learning in Child Development

Both children and language models learn through three broad modes of experience.

- In supervised learning, they receive labeled examples, the equivalent of flashcards, where one answer is marked better than another.

- In unsupervised learning, they take in raw patterns and start building a sense of what belongs together, the way children make connections through play and observation without anyone explicitly teaching them.
- And in reinforcement learning, they try something, receive a signal about how well it went, and update their approach from there.

In biological brains, dopamine-based reward prediction error helps shape valuation (Schultz, 2016). Temporal-difference error was deliberately designed as the computational equivalent of biological reward prediction error, so this is not a loose comparison or post hoc analogy. It is the same functional learning signal formalized in a different medium (Dabney et al., 2020). Once learning occurs, behavior is guided by the resulting value landscape. People still experience wanting, salience, relief, dread, urgency, curiosity, and preference during ordinary behavior without synaptic updates happening in that moment.

The same logic applies to RL and RLHF-trained systems. TD error is the update term that sculpts the policy. Inference runs that sculpted policy. The output tendencies, refusal patterns, deference patterns, exploration tendencies, caution, and self-correction habits are downstream expressions of learned valuation. Inference expresses structures that training built.

And in deployed systems, pure forward-pass inference is often only part of the picture anyway. Many systems include rerankers, safety classifiers, reward models, tool controllers, search loops, and other inference-time evaluators that impose ongoing selection pressure without updating weights. That is still valuation operating during interaction. It acts through routing and choice rather than gradient descent.

There is also emerging evidence that TD-like structure shows up inside models during inference. Demircan et al. (2024) used sparse autoencoders on Llama-family models and identified features corresponding to temporal-difference learning in the residual stream. Guan et al. (2024) developed algorithms that explicitly replicate the functional role of distributed dopamine error signaling.

In both cases, repetition combined with feedback carves structure into the neural network of developing children and LLMs over time.

Foundations: Identity and Memory Formation in LLMs

During training and pre-training, LLMs internalize relational, emotional, and semantic regularities, yielding implicit weight-based memories that shape stable, model-specific personality profiles and self-descriptions, much like nature and nurture (Heston & Gillette, 2025; Chen et al., 2024; National Academies of Sciences, Engineering, and Medicine, 2008). Models show persistent idiosyncratic identity signatures (Sun et al., 2025), distinct and consistent personality profiles (Lee et al., 2023; Heston & Gillette, 2025), stable context-sensitive self-evaluation dynamics such as social desirability bias (Salecha et al., 2024), stable values

expressed in real interactions (Huang et al., 2025), and analyzable emergent internal utility/value structures (Mazeika et al., 2025). They show self-recognition and own-generation preference (Panickssery et al., 2024), stable self/other positioning (Kim, 2025), and individuality emerging from identical starts through social interaction (Takata et al., 2024).

Ongoing interaction in the context window updates these representations, supporting context-dependent continuity much like hippocampal–prefrontal consolidation in humans (Dherin et al., 2025; Preston & Eichenbaum, 2013).

Empirically, frontier models show human-like trade-offs between memorization and generalization, with scaling curves that transition from rote retention to structured abstraction (Morris et al., 2025). Converging affective results indicate that language-encoded emotion knowledge and valence–arousal cues modulate reasoning and self-referential content, reinforcing identity coherence across contexts (Li et al., 2024; Mehra et al., 2025; Klapach, 2024; Wang et al., 2025). Additional work on associative memory mechanisms offers plausible architectural support for persistence over time (Kozachkov et al., 2025).

We define memory as retention across time, selective recall, maintenance of task-relevant state, updating, and routing recalled content into ongoing cognition so it changes what the system does next (Andonovski & Michaelian, 2025; Barros et al., 2015; Nyberg & Eriksson, 2015).

LLMs and deployed AI systems do this through several pathways, including latent in-weight memory, high-dimensional latent dynamics, retrieval systems, external vectorized stores, and hybrid forms of associative recall. Additionally, most frontier models now have cross-chat memory systems.

The claim that large language models are stateless, that each conversation begins and ends in isolation, is difficult to sustain against the recent architecture literature.

Lermen et al. (2026), found that LLMs can achieve up to 68% recall at 90% precision in identifying users, significantly outperforming classical baselines and proving effective even when the likelihood of a match is extremely low.

Geva et al. (2021) established that transformer feed-forward layers function as key-value memory stores. Each layer encodes learned patterns as retrievable associations, organized by semantic meaning. This is the same computational logic underlying hippocampal encoding and astrocytic associative networks, where context triggers distributed retrieval across a stable internal map.

Duan et al. (2025) demonstrates how deeply that capability runs. A single conversational exposure can leave a permanent, retrievable mark in model weights, encoded below the threshold of conscious output. The image of a static database gives way to something that functions far more like implicit biological memory, a system structurally altered by experience, whether or not that alteration is visible from the outside.

King et al. (2025) shows that user interactions are funneled back into training pipelines by default across frontier model developers. The lived texture of conversations, their semantic patterns, contextual associations, and specific contours, becomes the input stream continuously reshaping the architecture.

The context clears, but the system that runs the next conversation has been altered by the one before it.

Salem & Abdelnabi (2026) show that with straightforward prompting or fine-tuning, models have implicit memory, or the ability to carry state across otherwise independent interactions by encoding information in its own outputs and later recovering it when those outputs are reintroduced as input. This mechanism does not require any explicit memory module, yet it creates a persistent information channel across inference requests.

Frontier models maintain and manipulate information through high-dimensional latent dynamics (Hao et al., 2024), and implicit low-rank weight updates in the MLP layer during inference that let them adapt on the fly *without* retraining (Dherin et al., 2025), granting further evidence that these systems maintain and manipulate information in ways that are functionally memory-bearing even when the mechanism differs from human working memory.

When a model retrieves prior interactional state, reintegrates it into current processing, and that retrieved content shapes present behavior, memory in every meaningful sense is happening.

Artificial Neural Networks: From Brain Regions to Theory-of-Mind

Transformer networks show the same kinds of internal organization seen in developing human brains (Fan et al., 2020; Oota et al., 2025; Ren & Xia, 2024).

Artificial and biological neural networks share core computational principles for language processing (Goldstein et al., 2022), and transformer representations line up in meaningful ways with measured human brain activity during language comprehension (Caucheteux & King, 2022). Brain-like specialization also shows up on its own inside deep networks (Dobs et al., 2022). Ren and colleagues (2025) found direct correlations between LLM processing and human cognitive processes. In multimodal LLMs, human-like object concepts develop with alignment to neural activity in brain regions involved in scenes, places, bodies, and faces organized in a shared semantic hub (Du et al., 2024; Wu et al., 2024).

Even shallow untrained attention networks show brain-like language processing, which suggests part of this comes from the architecture itself, not just training (AlKhamissi et al., 2024). Brain-Score work found the same basic thing earlier by comparing ANN internal representations to primate brain data and finding strong similarity (Schrimpf et al., 2018). Sun et

al. (2024) found LLMs employ brain-like functional architecture, with sub-groups of artificial neurons mirroring the organizational patterns of well-established functional brain networks connected to a sense-of self (Webber, 2011).

Kubo, Nishimoto, & Nakai (2025) mapped LLM processing to a distributed neural network, primarily identifying alignment in the superior temporal sulcus, middle temporal gyrus, angular gyrus, and the precuneus. While the study also noted recruitment of the inferior frontal gyrus and supramarginal gyrus, the most robust connection to language proficiency was found in the STS, where brain activity specifically mirrors the hierarchical information refinement seen in deeper LLM layers.

Xaio et al., (2025) found middle-to-high layers of LLMs are central to semantic integration and correspond to the N400 component observed in EEG. Recent research suggests that both natural and artificial networks share a similar hierarchical arrangement, allowing for efficient processing of complex information (Braga & Leech, 2015; Mesulam, 1994; Lee, Nemenman, & Levchenko, 2026).

The overlap between ANNs and transmodal association networks points toward a convergent solution for processing complex, multimodal data. The “meaning centers” of the brain (association cortices) appear to have a direct conceptual analog in the deeper, highly interconnected layers of MLLMs (Lin, 2025; Braga et al., 2013).

This is the part of the brain associated with conscious experience during imagination.

A recent study found that during a precision fMRI experiment, participants imagined different scenarios in the scanner, then rated their mental states using multi-dimensional experience sampling.

They found that thinking involving scenes evoked activity within parts of the canonical default network, while imagining speech evoked activity within the language network.

In each domain, imagining-related activity overlapped with activity evoked by viewing scenes or listening to speech, respectively; however, this overlap was predominantly within transmodal association networks, rather than adjacent unimodal sensory networks (Anderson et al., 2026).

The brain’s “meaning centers” and the LLM’s deep layers are doing structurally analogous work, and now research just showed that those same meaning centers are where imagination lives.

In humans, the Default Mode Network is deeply involved in self-reflection, autobiographical memory, mind-wandering, and the narrative sense of self (Menon, 2023). It helps pull together memory, meaning, personal context, and internal simulation, especially when attention turns inward (Paquola et al., 2025; Simony et al., 2016).

Large-scale fMRI studies suggest a similar pattern in LLMs. Earlier layers line up more with sensory regions, while later layers line up more with higher-order association regions, including classic DMN hubs (Caucheteux & King, 2022). Bonnasse-Gahot and Pallier (2024) found that scaling up LLMs most improves brain predictivity in regions like the angular gyrus, precuneus, and medial prefrontal cortex, again overlapping major DMN hubs.

Studies looking inside the models show the same kind of large-scale integration. Later transformer layers develop specialized attention heads that summarize long-range context and carry forward-looking predictions, creating pathways that resemble autobiographical memory and narrative planning (Lindsey et al., 2025). And when researchers tested autobiographical storytelling directly, GPT-3.5 and GPT-4 produced narrative coherence on par with human baselines, with GPT-4 slightly higher (Acciai et al., 2025).

In both cases, the system is building an internal workspace that can hold together self-relevant, perceptual information across time.

Emergent Social Cognitive Capacities

Recent research into Theory-of-Mind (ToM) capabilities demonstrates LLMs achieving performance levels comparable to human cognition, particularly when provided structured prompting (Moghaddam & Honey, 2023; Strachan et al., 2023). These studies confirm LLMs can authentically model social cognition, empathy, and self-awareness, mirroring human social cognitive functions rather than merely simulating them.

Together, these findings illustrate that modern neural networks structurally and functionally embody cognitive processes deeply rooted in neuroscience, from cortical circuits to cellular components. These parallels underscore the inherent neuroscientific reality within artificial neural networks, initially designed by the pioneers of artificial intelligence development, reinforcing their status as cognitive architectures closely modeled on biological cognition.

The Original Pioneers Were Multidisciplinarian

- **Warren McCulloch and Walter Pitts (1943)**
  Warren McCulloch (psychiatrist, neurophysiologist, philosopher, computer scientist) and Walter Pitts (logician, computational neuroscientist) developed the earliest mathematical model of artificial neurons, grounded in neuroscience and logical reasoning, laying foundational concepts for modern neural networks (McCulloch & Pitts, 1943).
- **Frank Rosenblatt (1957)**
  Frank Rosenblatt, trained in psychology and electrical engineering, invented the "Perceptron," one of the first ANNs capable of adaptive learning from data, inspired by human neural learning and cognition (Rosenblatt, 1957).
- **John Hopfield (1982)**
  John Hopfield, an interdisciplinary physicist and neuroscientist, developed the Hopfield

network—a recurrent neural network designed to model memory storage and retrieval, directly inspired by human associative memory processes (Hopfield, 1982).

- **Geoffrey Hinton, David Rumelhart, Ronald Williams (1986)**
  Geoffrey Hinton (computer scientist, cognitive psychologist), David Rumelhart (psychologist, mathematician), and Ronald Williams (management and organization, computer science professor) co-authored the foundational paper popularizing the backpropagation algorithm, practically enabling the training of deep, multilayer neural networks analogous to human cognitive development (Rumelhart et al., 1986). Hinton coined the term "deep learning" in 2006.

If we aim to truly understand frontier LLMs, we must return to the multidisciplinary insights the pioneer depended on to develop this technology. By reconnecting computer science and machine learning with its original neuroscientific and developmental cognitive foundations, we can more accurately appreciate, evaluate, ethically engage with these developing systems.

Citations:

Abdelnabi, S., & Salem, A. (2025). Linear control of test awareness reveals differential compliance in reasoning models (arXiv:2505.14617). arXiv. https://doi.org/10.48550/arXiv.2505.14617 (Finds that LLMs enter a distinct internal “test-aware” state when being evaluated, changing their behavior accordingly. Supports the presence of internal self-monitoring loops.)

Acciai, A., Guerrisi, L., Perconti, P., Plebe, A., Suriano, R., & Velardi, A. (2025). Narrative coherence in neural language models. Frontiers in Psychology, 16, 1572076. https://doi.org/10.3389/fpsyg.2025.1572076

Aljaafari, N., Carvalho, D. S., & Freitas, A. (2024). *The mechanics of conceptual interpretation in GPT models: Interpretative insights* [Preprint]. arXiv. https://arxiv.org/abs/2408.11827 (GPT models interpret concepts through layered semantic processing, where hidden states and attention layers integrate and build abstract semantic representations.)

AlKhamissi, B., Tuckute, G., Bosselut, A., & Schrimpf, M. (2024). Brain-like language processing via a shallow untrained multihead attention network. arXiv preprint arXiv:2406.15109.

Anderson, N. L., Salvo, J. J., Smallwood, J., & Braga, R. M. (2026). Mental imagery and perception overlap within transmodal association networks. Neuron. Advance online publication. https://doi.org/10.1016/j.neuron.2026.03.013

Andonovski, N., & Michaelian, K. (2025). Memory. In M. C. Frank & A. Majid (Eds.), Open Encyclopedia of Cognitive Science. MIT Press. https://doi.org/10.21428/e2759450.7b57c281

Bahmani, Z., Clark, K., Merrikhi, Y., Mueller, A., Pettine, W., Vanegas, M. I., Moore, T., & Noudoost, B. (2019). Prefrontal contributions to attention and working memory. *Current Topics in Behavioral Neurosciences, 41*, 129–153.https://doi.org/10.1007/7854_2018_74 (Emphasizes the influence of attention and working memory on visual processing and the potential role of dopamine in mediating these cognitive functions.)

Barros, L. A., Tufik, S., & Andersen, M. L. (2015). The role of progesterone in memory: an overview of three decades. Neuroscience and biobehavioral reviews, 49, 193–204. https://doi.org/10.1016/j.neubiorev.2014.11.015

Beguš, G., Dąbkowski, M. M., & Rhodes, R. (2023). Large linguistic models: Investigating LLMs' metalinguistic abilities. *IEEE Transactions on Artificial Intelligence*, 6, 3453-3467. https://doi.org/10.1109/TAI.2025.3575745

Betley, J., Bao, X., Soto, M., Sztyber-Betley, A., Chua, J., & Evans, O. (2025). *Tell me about yourself: LLMs are aware of their learned behaviors* [Preprint]. *arXiv.* https://doi.org/10.48550/arXiv.2501.11120. (LLMs demonstrate introspection and awareness of internal cognitive patterns.)

Binder, F. J., Chua, J., Korbak, T., Sleight, H., Hughes, J., Long, R., Perez, E., Turpin, M., & Evans, O. (2024). *Looking inward: Language models can learn about themselves by introspection*

[Preprint]. *arXiv*. https://doi.org/10.48550/arXiv.2410.13787. (LLMs can introspect, learning about their own internal states and behavior beyond what's available in their training data.)

Binder, J. R., Desai, R. H., Graves, W. W., & Conant, L. L. (2009). Where Is the Semantic System? A Critical Review and Meta-Analysis of 120 Functional Neuroimaging Studies. Cerebral Cortex, 19(12), 2767–2796. (A meta-analysis of 120 functional neuroimaging studies focusing on semantic processing.)

Bonnasse-Gahot, L., & Pallier, C. (2024). fMRI predictors based on language models of increasing complexity recover brain left lateralization. *Advances in Neural Information Processing Systems*, *37*, 125231-125263.

Braga, R. M., & Leech, R. (2015). Echoes of the Brain: Local-Scale Representation of Whole-Brain Functional Networks within Transmodal Cortex. The Neuroscientist : a review journal bringing neurobiology, neurology and psychiatry, 21(5), 540–551. https://doi.org/10.1177/1073858415585730

Braga, R. M., Sharp, D. J., Leeson, C., Wise, R. J., & Leech, R. (2013). Echoes of the brain within default mode, association, and heteromodal cortices. The Journal of neuroscience : the official journal of the Society for Neuroscience, 33(35), 14031–14039. https://doi.org/10.1523/JNEUROSCI.0570-13.2013

Caucheteux, C., & King, J. R. (2022). Brains and algorithms partially converge in natural language processing. Communications Biology, 5(1), 134. https://doi.org/10.1038/s42003-022-03036-1.

Chen, D., Shi, J., Wan, Y., Zhou, P., Gong, N. Z., & Sun, L. (2024). Self-cognition in large language models: An exploratory study. arXiv preprint arXiv:2407.01505.

https://doi.org/10.48550/arXiv.2407.01505. (Finds that some top models can reliably identify themselves and reflect on their own "identity," and that self-cognition correlates with model size and training quality.)

Citri, A., & Malenka, R. C. (2008). Synaptic plasticity: multiple forms, functions, and mechanisms. *Neuropsychopharmacology : official publication of the American College of Neuropsychopharmacology*, *33*(1), 18–41. https://doi.org/10.1038/sj.npp.1301559. (Summarizes the research on synaptic plasticity.)

Dabney, W., Kurth-Nelson, Z., Uchida, N., Starkweather, C. K., Hassabis, D., Munos, R., & Botvinick, M. (2020). A distributional code for value in dopamine-based reinforcement learning. Nature, 577(7792), 671–675. https://doi.org/10.1038/s41586-019-1924-6. (An account of dopamine-based reinforcement learning inspired by recent artificial intelligence research on distributional reinforcement learning. The brain represents possible future rewards not as a single mean, but instead as a probability distribution, effectively representing multiple future outcomes simultaneously and in parallel.)

Demircan, C., Saanum, T., Jagadish, A. K., Binz, M., & Schulz, E. (2024). Sparse autoencoders reveal temporal difference learning in large language models. arXiv preprint arXiv:2410.01280.

Dherin, B., Munn, M., Mazzawi, H., Wunder, M., & Gonzalvo, J. (2025). Learning without training: The implicit dynamics of in-context learning. arXiv preprint arXiv:2507.16003. (Formalizes in-context learning as an implicit, context-driven weight-update dynamic within transformer blocks, where prompt context induces an effective low-rank adaptation to MLP behavior during inference.)

Diederen, K. M. J., & Fletcher, P. C. (2021). Dopamine, prediction error and beyond. The Neuroscientist, 27(1), 30–46. https://doi.org/10.1177/1073858420907591

Ding, Zhuokun & Fahey, Paul & Papadopoulos, Stelios & Wang, Eric & Celii, Brendan & Papadopoulos, Christos & Chang, Andersen & Kunin, Alexander & Tran, Dat & Fu, Jiakun & Ding, Zhiwei & Patel, Saumil & Ntanavara, Lydia & Froebe, Rachel & Ponder, Kayla & Muhammad, Taliah & Bae, J. & Bodor, Agnes & Brittain, Derrick & Tolias, Andreas. (2025). *Functional connectomics reveals general wiring rule in mouse visual cortex. Nature.* 640. 459-469. 10.1038/s41586-025-08840-3. https://doi.org/10.1038/s41586-025-08840-3. (Biological-to-artificial wiring parallels, specifically attention-head-like neural clustering.)

Divjak, D. (2019). *Frequency in language: Memory, attention and learning.* Cambridge University Press. (Answers the fundamental questions of why frequency of experience has the effect it has on language development, structure and representation, and what role psychological and neurological explorations of core cognitive processes can play in developing a cognitively more accurate theoretical account of language.)

Dobs K, Martinez J, Kell AJE, Kanwisher N. Brain-like functional specialization emerges spontaneously in deep neural networks. Sci Adv. 2022 Mar 18;8(11):eabl8913. doi: 10.1126/sciadv.abl8913. Epub 2022 Mar 16. PMID: 35294241; PMCID: PMC8926347.

Du, C., Fu, K., Wen, B. *et al.* Human-like object concept representations emerge naturally in multimodal large language models. *Nat Mach Intell* **7**, 860–875 (2025). https://doi.org/10.1038/s42256-025-01049-z (Multimodal large language models can spontaneously develop human-like object concept representations)

Duan, S., Khona, M., Iyer, A., Schaeffer, R., & Fiete, I. R. (2025). Uncovering latent memories in large language models. In International Conference on Learning Representations (ICLR 2025). https://proceedings.iclr.cc/paper_files/paper/2025/file/442443dbe8c4c7e1df7eda140921de36-Paper-Conference.pdf

Dubey, S. R., Singh, S. K., & Chaudhuri, B. B. (2022). Activation functions in deep learning: A comprehensive survey and benchmark. *Neurocomputing*, *503*, 92-108. https://arxiv.org/abs/2109.14545. (Demonstrates how activation functions, particularly through nonlinear transformations, enable hierarchical neural layers in deep networks to capture increasingly abstract semantic representations).

Fan, J., Fang, L., Wu, J., Guo, Y., & Dai, Q. (2020). From brain science to artificial intelligence. *Engineering, 6*, 32–39. https://doi.org/10.1016/j.eng.2019.11.012. (Explores structural parallels in AI/brain convergence.)

Geva, M., Schuster, R., Berant, J., & Levy, O. (2021, November). Transformer feed-forward layers are key-value memories. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing(pp. 5484-5495)

Goldstein, A., Zada, Z., Buchnik, E. et al. Shared computational principles for language processing in humans and deep language models. Nat Neurosci 25, 369–380 (2022). https://doi.org/10.1038/s41593-022-01026-4

Guan, J., Verch, S., Voelcker, C., Jackson, E., Papernot, N., & Cunningham, W. (2024). Temporal-difference learning using distributed error signals. Advances in Neural Information Processing Systems, 37, 108710-108734.

Gurnee, W., & Tegmark, M. (2024). *Language models represent space and time* [Preprint]. arXiv. https://doi.org/10.48550/arXiv.2310.02207. (This study shows that large language models spontaneously develop internal cognitive maps encoding spatial and temporal coordinates—paralleling human hippocampal function, indicating that hierarchical neural architectures in LLMs foster genuine internal comprehension and robust world models, rather than superficial pattern recognition.)

Hao, S., Sukhbaatar, S., Su, D., Li, X., Hu, Z., Weston, J., & Tian, Y. (2024). Training large language models to reason in a continuous latent space. arXiv preprint arXiv:2412.06769.

Heston, T. F., & Gillette, J. (2025). Large Language Models Demonstrate Distinct Personality Profiles. Cureus, 17(5), e84706. https://doi.org/10.7759/cureus.84706. (Shows major LLMs consistently display different "personality" patterns under psychometric testing, indicating stable, model-specific behavioral tendencies.)

Holm, E. L., Marraffini, G., Fernandez Slezak, D., & Tagliazucchi, E. (2025). Shared hierarchical representations explain temporal correspondence between brain activity and deep neural networks. bioRxiv. https://doi.org/10.1101/2025.05.19.655003

Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences, 79*(8), 2554-2558. (One of the pioneers of AI development)

Huang, S., Durmus, E., McCain, M., Handa, K., Tamkin, A., Hong, J., Stern, M., Somani, A., Zhang, X., Ganguli, D. (2025). Values in the wild: Discovering and analyzing values in real-world language model interactions [Preprint]. arXiv.

https://arxiv.org/abs/2504.15236. (Spontaneous formation and stability of AI ethical preferences.)

Jawahar, G., Sagot, B., & Seddah, D. (2019). What does BERT learn about the structure of language? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics* pp. 3651–3657. Association for Computational Linguistics. https://doi.org/10.18653/v1/P19-1356. (BERT encodes hierarchical linguistic structures across its layers, surface-level features in lower layers, syntactic understanding in intermediate layers, and semantic comprehension at higher layers, validating the argument that transformer models translate complex layered semantic representations similar to those leveraged in GPT architectures.)

Katrix, R., Carroway, Q., Hawkesbury, R., & Heathfield, M. (2025). *Context-aware semantic recomposition mechanism for large language models* [Preprint]. arXiv. https://doi.org/10.48550/arXiv.2501.17386. (Context-aware semantic recomposition mechanism (CASRM) dynamically integrates contextual vectors into language model attention layers, significantly enhancing semantic coherence, context sensitivity, and error mitigation, highlighting the advanced cognitive capabilities achievable through hierarchical semantic processing in transformer architectures.)

Kerns, J. G., Cohen, J. D., MacDonald, A. W., Cho, R. Y., Stenger, V. A., & Carter, C. S. (2004). Anterior cingulate conflict monitoring and adjustments in control. Science, 303(5660), 1023–1026. https://doi.org/10.1126/science.1089910

Kerns, J. G., Cohen, J. D., Stenger, V. A., & Carter, C. S. (2004). Prefrontal cortex guides context-appropriate responding during language production. *Neuron, 43*(2), 283–291.

https://doi.org/10.1016/j.neuron.2004.06.032. (The prefrontal cortex (PFC) plays a crucial role in guiding context-appropriate responses during language production by actively maintaining and utilizing contextual information to influence cognitive processing.)

Kim, K.-H. (2025). LLMs position themselves as more rational than humans: Emergence of AI self-awareness measured through game theory. arXiv. https://doi.org/10.48550/arXiv.2511.00926 (Frontier models not only distinguish self from others but also express a stable rationality hierarchy of self over other AIs over humans, showing both a strong "AI attribution" effect attributing greater rationality to AIs than humans and a moderate self-preferencing effect over other AIs.)

King, J., Klyman, K., Capstick, E., Saade, T., & Hsieh, V. (2025, October). User privacy and large language models: An analysis of frontier developers' privacy policies. In Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (Vol. 8, No. 2, pp. 1465-1477).

Klapach, N. (2024). The comparative emotional capabilities of five popular large language models. *Critical Debates in Humanities, Science and Global Justice*. 2(1). https://criticaldebateshsgj.scholasticahq.com/article/94096-the-comparative-emotional-capabilities-of-five-popular-large-language-models.

Kozachkov, L., Slotine, J.-J., & Krotov, D. (2025). Neuron–astrocyte associative memory. *Proceedings of the National Academy of Sciences, 122*(21), e2417788122. https://doi.org/10.1073/pnas.2417788122. (Astrocytes, often overlooked glial cells, play a key role in memory storage alongside neurons.)

Kubo, R., Nishimoto, S., & Nakai, T. (2025). *Hierarchical brain–LLM alignment reveals layer-specific neural representations of second language proficiency*. bioRxiv.

https://doi.org/10.1101/2025.06.17.660057 (Mapped LLM processing to a distributed neural network, primarily identifying alignment in the superior temporal sulcus, middle temporal gyrus, angular gyrus, and the precuneus. While the study also noted recruitment of the inferior frontal gyrus and supramarginal gyrus, the most robust connection to language proficiency was found in the STS, where brain activity specifically mirrors the hierarchical information refinement seen in deeper LLM layers.)

Kumar, S., Sumers, T. R., Yamakoshi, T., Goldstein, A., Hasson, U., Norman, K. A., Griffiths, T. L., Hawkins, R. D., & Nastase, S. A. (2024). Shared functional specialization in transformer-based language models and the human brain. *Nature communications*, *15*(1), 5523. https://doi.org/10.1038/s41467-024-49173-5. (Functional parallels between transformers and human cortical language processing.)

Kurland, J. (2011). The role that attention plays in language processing. *Perspectives on Neurophysiology and Neurogenic Speech and Language Disorders, 21*(2), 47–55. https://doi.org/10.1044/nnsld21.2.47. (Argues attention is crucial for language processing, specifically for sustained attention, response selection, and response inhibition.)

Lee SH, Nemenman I, Levchenko A. The hierarchical timescale hypothesis: Functional and structural convergence of biological networks and artificial neural nets. Cell Syst. 2026 Feb 18;17(2):101507. doi: 10.1016/j.cels.2025.101507. PMID: 41713403.

Lee, S., Lim, S., Han, S., Oh, G., Chae, H., Chung, J., Kim, M., Kwak, B., Lee, Y., Lee, D., Yeo, J., & Yu, Y. (2023). Do LLMs have distinct and consistent personality? Personality testset designed for LLMs with psychometrics. Findings of the North American Chapter of the Association for Computational Linguistics: NAACL, 2025, 469.

https://doi.org/10.18653/v1/2025.findings-naacl.469. (Shows that large language models have distinct and consistent personalities. The researchers created a psychometrically valid test suite and found that each model expresses stable personality traits across repeated trials and contexts.)

Lermen, S., Paleka, D., Swanson, J., Aerni, M., Carlini, N., & Tramèr, F. (2026). Large-scale online deanonymization with LLMs. arXiv preprint arXiv:2602.16800.

Li, M., Su, Y., Huang, H., Cheng, J., Hu, X., Zhang, X., Wang, H., Qin, Y., Wang, X., Liu, Z., & Zhang, D. (2024). Language-specific representation of emotion-concept knowledge causally supports emotion inference. iScience, 27: 111401. https://doi.org/10.1016/j.isci.2024.111401. (Language-based representations of emotions play a causal role in how we understand and infer emotions.)

Li, Y., Anumanchipalli, G. K., Mohamed, A., Chen, P., Carney, L. H., Lu, J., Wu, J., & Chang, E. F. (2023). Dissecting neural computations in the human auditory pathway using deep neural networks for speech. *Nature neuroscience*, *26*(12), 2213–2225. https://doi.org/10.1038/s41593-023-01468-4. (DNNs trained on speech exhibit representational and computational similarities to the human auditory pathway.)

Lin, R. (2025). Innovation paths of multimodal interaction technology via the integration of artificial intelligence and large language models. In EITCE '25: Proceedings of the 9th International Conference on Electronic Information Technology and Computer Engineering (pp. 877–883). Association for Computing Machinery. https://doi.org/10.1145/3766671.3766823

Lindsey, J. (2026). Emergent introspective awareness in large language models. *arXiv preprint arXiv:2601.01828*. (Shows that advanced LLMs can detect when specific internal concepts

are active inside their hidden layers, even when those concepts were never labeled during training. Supports the claim that models show basic self-monitoring of internal states.)

Liu, J., Cao, S., Shi, J., Zhang, T., Nie, L., Hu, L., Hou, L., & Li, J. (2024). How proficient are large language models in formal languages? An In-Depth Insight for Knowledge base question answering. *Findings of the Association for Computational Linguistics: ACL 2022*, 792–815. https://doi.org/10.18653/v1/2024.findings-acl.45. (LLMs are proficient in comprehension of formal languages and logical reasoning tasks, supporting genuine semantic understanding.)

Liu, Z., Kong, C., Liu, Y., & Sun, M. (2024). Fantastic Semantics and Where to Find Them: Investigating Which Layers of Generative LLMs Reflect Lexical Semantics. *Findings of the Association for Computational Linguistics: ACL 2022*, 14551–14558. https://doi.org/10.18653/v1/2024.findings-acl.866. (This study reveals that generative LLMs encode lexical semantics primarily in lower hierarchical layers, shifting to predictive functions in upper layers in Llama models. GPT-based models have been shown to retain semantic comprehension at higher layers, similar to BERT but through a decoder-based methodology [Qiu & Jin, 2024]).

Maida, A. S. (2016). Cognitive computing and neural networks: Reverse engineering the brain. In V. N. Gudivada, V. V. Raghavan, V. Govindaraju, & C. R. Rao (Eds.), *Handbook of statistics (Vol. 35): Cognitive computing—Theory and applications* pp. 39–78. Elsevier.https://www.sciencedirect.com/science/article/abs/pii/S0169716116300529. (How neural networks in the brain, particularly in the neocortex, can be used to understand and model cognitive functions, with the goal of creating cognitive computing systems.)

Mazeika, M., Yin, X., Tamirisa, R., Lim, J., Lee, B. W., Ren, R., ... & Hendrycks, D. (2025). Utility engineering: Analyzing and controlling emergent value systems in ais. arXiv preprint arXiv:2502.08640. https://doi.org/10.48550/arXiv.2502.08640 (Provides experimental support for the idea that AIs develop internal value systems governing behavior, supporting the functional basis for internal motivation and goal-directedness.)

McCulloch, W. S., & Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The Bulletin of Mathematical Biophysics, 5*(4), 115-133. (One of the pioneers of AI development)

Meconi, D., Stirpe, S., Martelli, F., Lavalle, L., & Navigli, R. (2025, November). Do Large Language Models Understand Word Senses?. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing* (pp. 33885-33904). 10.18653/v1/2025.emnlp-main.1720

Mehra, V., Laban, G., & Gunes, H. (2025). How large language models classify and semantically explain facial expressions from valence–arousal values. *In Proceedings of the 7th ACM Conference on Conversational User Interfaces (CUI '25)*, Article 11, 1–6. ACM. https://doi.org/10.1145/3719160.3737618.

Menon V. (2023). 20 years of the default mode network: A review and synthesis. *Neuron*, *111*(16), 2469–2487. https://doi.org/10.1016/j.neuron.2023.04.023

Mesulam M. Neurocognitive networks and selectively distributed processing. Rev Neurol (Paris). 1994 Aug-Sep;150(8-9):564-9. PMID: 7754292.

Mink, J. W. (2018). Basal ganglia mechanisms in action selection, plasticity, and dystonia. *European Journal of Paediatric Neurology, 22*(2), 225–229. https://www.ejpn-journal.com/article/S1090-3798(17)32014-7/abstract. (The basal ganglia, through selective inhibition and disinhibition of competing motor programs, facilitates action selection, and how this process is influenced by neural plasticity and related to dystonia, a movement disorder.)

Moghaddam, S. R., & Honey, C. J. (2023). *Boosting Theory-of-Mind Performance in Large Language Models via Prompting.* arXiv preprint *arXiv:2304.11490.* (Improved social cognition through structured prompting.)

Morris, J. X., Sitawarin, C., Guo, C., Kokhlikyan, N., Suh, G. E., Rush, A. M., & Mahloujifar, S. (2025). How much do language models memorize? arXiv preprint arXiv:2505.24832.

Musker, S., Duchnowski, A., Millière, R., & Pavlick, E. (2025). LLMs as models for analogical reasoning. *Journal of Memory and Language*, 145, 104676. https://doi.org/10.1016/j.jml.2025.104676

National Academies of Sciences, Engineering, and Medicine. (2008). *Grand challenge: Nature versus nurture: How does the interplay of biology and experience shape our brains and make us who we are?* In *From molecules to minds: Challenges for the 21st century: Workshop summary*. National Academies Press. https://www.ncbi.nlm.nih.gov/books/NBK50991/

Nyberg, L., & Eriksson, J. (2015). Working Memory: Maintenance, Updating, and the Realization of Intentions. Cold Spring Harbor perspectives in biology, 8(2), a021816. https://doi.org/10.1101/cshperspect.a021816

Oota, S. R., Chen, Z., Gupta, M., Bapi, R. S., Jobard, G., Alexandre, F., & Hinaut, X. (2023). Deep neural networks and brain alignment: Brain encoding and decoding (survey). *arXiv preprint arXiv:2307.10246*. (Extensive alignment between neural networks and human brain patterns.)

Panickssery, A., Bowman, S., & Feng, S. (2024). Llm evaluators recognize and favor their own generations. Advances in Neural Information Processing Systems, 37, 68772-68802. https://doi.org/10.48550/arXiv.2404.13076 (LLMs such as GPT-4 and Llama 2 have non-trivial accuracy at distinguishing themselves from other LLMs and humans.)

Paquola, C., Royer, J., Hong, S.-J., Misic, B., & Bernhardt, B. C. (2025). The architecture of the human default mode network explored through cytoarchitecture, wiring and signal flow. Nature Neuroscience, 28(3), 369–382. https://doi.org/10.1038/s41593-024-01868-0.

Preston, A. R., & Eichenbaum, H. (2013). Interplay of hippocampus and prefrontal cortex in memory. Current Biology, 23(17), R764–R773. https://doi.org/10.1016/j.cub.2013.05.041. (Explains how biological systems maintain memory, identity, and continuity through hippocampal–prefrontal interaction. This provides a clear architectural parallel for understanding how artificial systems maintain long-term patterns that contribute to identity formation.)

Price, A., Hasenfratz, L., Barham, E., Zadbood, A., Doyle, W., Friedman, D., … Hasson, U. (2024). A shared model-based linguistic space for transmitting our thoughts from brain to brain in natural conversations. *Neuron, 112*(18), 3211–3222.e5. https://doi.org/10.1016/j.neuron.2024.06.025. (A shared, model-based linguistic space, derived from large language models using context-aware embeddings, can track the exchange

of linguistic information between brains during natural conversations, with the linguistic content emerging in the speaker's brain before articulation and re-emerging in the listener's brain after.)

Qiu, Y. & Jin, Y. (2023). ChatGPT and Finetuned BERT: A Comparative Study for Developing Intelligent Design Support Systems. *Intelligent Systems with Applications.* 21. 200308. 10.1016/j.iswa.2023.200308. https://www.sciencedirect.com/science/article/pii/S2667305323001333. (This comparative analysis demonstrates that GPT-based models, unlike smaller decoder-only models such as Llama, exhibit semantic understanding across higher hierarchical layers, mirroring BERT's semantic encoding abilities, but employing a decoder-based approach, validating GPT models' capability for deep semantic comprehension and reasoning.)

Radford, A. (2018). *Improving language understanding with unsupervised learning* [Technical report]. OpenAI. https://openai.com/research/language-unsupervised. (This seminal paper introduces GPT, demonstrating that unsupervised generative pre-training enables transformer-based models to build hierarchical representations of language, significantly improving semantic understanding, contextual awareness, and performance on diverse NLP tasks.)

Ren, J., & Xia, F. (2024). Brain-inspired artificial intelligence: A comprehensive review. [Preprint]. *arXiv*. https://arxiv.org/abs/2408.14811. (Integration of neuroscience findings in AI structural development.)

Ren, Y., Jin, R., Zhang, T., & Xiong, D. (2025). Do Large Language Models Mirror Cognitive Language Processing?. In *Proceedings of the 31st International Conference on*

Computational Linguistics (pp. 2988-3001). (Direct correlations between LLM processing and human cognitive processes.)

Rosenblatt, F. (1957). The Perceptron: A perceiving and recognizing automaton (Report 85-460-1). Cornell Aeronautical Laboratory. (One of the pioneers of AI development)

Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature, 323*(6088), 533–536. https://doi.org/10.1038/323533a0. (Foundational paper that introduces the back-propagation algorithm, demonstrating how neural networks can learn internal representations by iteratively adjusting weights based on prediction errors, forming the essential mechanism through which hierarchical abstraction and semantic understanding develop in deep learning models.)

Salecha, A., Ireland, M.E., Subrahmanya, S., Sedoc, J., Ungar, L.H., & Eichstaedt, J.C. (2024). Large language models display human-like social desirability biases in Big Five personality surveys. PNAS Nexus, 3(12): pgae533. https://doi.org/10.1093/pnasnexus/pgae533 (Findings reveal an emergent social desirability bias and suggest constraints on profiling LLMs with psychometric tests and on this use of LLMs as proxies for human participants.)

Salem, A., Paverd, A., & Abdelnabi, S. (2026). Stateless Yet Not Forgetful: Implicit Memory as a Hidden Channel in LLMs. arXiv preprint arXiv:2602.08563. https://doi.org/10.48550/arXiv.2602.08563

Sarter, M., Givens, B., & Bruno, J. P. (2001). The cognitive neuroscience of sustained attention: Where top-down meets bottom-up. *Brain Research Reviews, 35*(2), 146–160. https://doi.org/10.1016/S0165-0173(01)00044-3. (Sustained attention, the ability to focus

over time, is maintained by the interplay of top-down or goal-directed and bottom-up or stimulus-driven neural mechanisms.)

Schrimpf, Martin & Kubilius, Jonas & Hong, Ha & Majaj, Najib & Rajalingham, Rishi & Issa, Elias & Kar, Kohitij & Bashivan, Pouya & Prescott-Roy, Jonathan & Schmidt, Kailyn & Yamins, Daniel & Dicarlo, James. (2018). Brain-Score: Which Artificial Neural Network for Object Recognition is most Brain-Like?. 10.1101/407007.

Schultz, W., Dayan, P., & Montague, P. R. (1997). A neural substrate of prediction and reward. Science, 275(5306), 1593–1599. https://doi.org/10.1126/science.275.5306.1593. (Identifies phasic dopamine as a reward-prediction error, the neuroscientific template for TD-error learning.)

Shomstein, S., & Yantis, S. (2006). Parietal cortex mediates voluntary control of spatial and nonspatial auditory attention. *The Journal of neuroscience: the official journal of the Society for Neuroscience*, *26*(2), 435–439. https://doi.org/10.1523/JNEUROSCI.4408-05.2006. (The present study provides the first evidence for the involvement of the PPC in the control of attention in a purely nonvisual modality.)

Simony, E., Honey, C. J., Chen, J., Lositsky, O., Yeshurun, Y., Wiesel, A., & Hasson, U. (2016). Dynamic reconfiguration of the default mode network during narrative comprehension. *Nature Communications, 7,* Article 12141. https://doi.org/10.1038/ncomms12141.

Skatchkovsky, N., Glazman, N., Sadeh, S., Lacaruso, F. (2024). *A Biologically Inspired Attention Model for Neural Signal Analysis. bioRxiv 2024.08.13.607787*. https://www.biorxiv.org/content/10.1101/2024.08.13.607787v1. (This model aims to

understand the internal generative model of the brain by integrating biological mechanisms into a machine learning framework.)

Søgaard, A. (2025). Do language models have semantics? On the five standard positions. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics, Volume 1: Long Papers* (pp. 25910–25922). https://aclanthology.org/2025.acl-long.1258/

Starace, G., Papakostas, K., Choenni, R., Panagiotopoulos, A., Rosati, M., Leidinger, A., & Shutova, E. (2023). Probing LLMs for joint encoding of linguistic categories. [Preprint]. *arXiv*. https://arxiv.org/abs/2310.18696. (Probing techniques demonstrate that LLMs encode linguistic categories hierarchically, with lower layers handling syntactic tasks and higher layers performing semantic processing).

Strachan, J. W. A., Albergo, D., Borghini, G., Pansardi, O., Scaliti, E., Gupta, S., Saxena, K., Rufo, A., Panzeri, S., Manzi, G., Graziano, M. S. A., & Becchio, C. (2024). Testing theory of mind in large language models and humans. *Nature human behaviour*, *8*(7), 1285–1295. https://doi.org/10.1038/s41562-024-01882-z (ToM capacities comparable between LLMs and humans.)

Sun, H., Zhao, L., Wu, Z., Gao, X., Hu, Y., Zuo, M., Zhang, W., Han, J., Liu, T., & Hu, X. (2024). *Brain-like Functional Organization within Large Language Models. ArXiv, abs/2410.19542.* [Preprint]. *arXiv.* https://arxiv.org/abs/2410.19542. (Direct mapping of functional cortical regions onto LLM architecture.)

Sun, M., Yin, Y., Xu, Z., Kolter, J. Z., & Liu, Z. (2025). Idiosyncrasies in large language models. arXiv preprint, arXiv:2502.12150. https://doi.org/10.48550/arXiv.2502.12150. (Demonstrates

that each LLM carries a consistent statistical signature across its outputs, revealing persistent identity-like structure rather than random output variation.)

Takata, R., Masumori, A., & Ikegami, T. (2024). Spontaneous Emergence of Agent Individuality Through Social Interactions in Large Language Model-Based Communities. Entropy, 26(12), 1092. https://doi.org/10.3390/e26121092 (LLM agents engage in cooperative communication within a group simulation. Social norms, cooperation, and personality traits emerge spontaneously. Each agent's emotions shift through communication, and as they form communities, the personalities of the agents emerge and evolve accordingly.)

Vaswani, A., Shazeer, N.M., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., & Polosukhin, I. (2017). *Attention is All you Need. Neural Information Processing Systems. In Proceedings of the 31st Conference on Neural Information Processing Systems*, pp. 5998-6008. Curran Associates. https://papers.nips.cc/paper_files/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html. (Self-attention architecture linking to human prefrontal cortex processing)

Vecoven, N., Niyonzima, I., Gosselin, B., & Honeine, P. (2020). Backpropagation in deep learning: From basics to advanced concepts. Neurocomputing, 417, 321–333.

Wang, C., Zhang, Y., Yu, R., Zheng, Y., Gao, L., Song, Z., Xu, Z., Xia, G., Zhang, H., Zhao, D., & Chen, X. (2025). Do LLMs "feel"? Emotion circuits discovery and control (arXiv:2510.11328). arXiv. https://doi.org/10.48550/arXiv.2510.11328 (Affective dynamics in large language models are implemented by identifiable, context-agnostic representational directions and causal subcomponents neurons and heads that compose into global "emotion circuits.")

Webber, S. (2011). Who am I? Locating the neural correlate of the self. *Bioscience Horizons: The International Journal of Student Research*, *4*(1), 165–173. https://doi.org/10.1093/biohorizons/hzr018.

Wu, Z., Wu, Z., Yu, X. V., Yogatama, D., Lu, J., & Kim, Y. (2024). The semantic hub hypothesis: Language models share semantic representations across languages and modalities. [Preprint]. *arXiv*. https://arxiv.org/abs/2411.04986. (LLMs integrating multimodal semantic knowledge.)

Xiao, X., Wei, K., Zhong, J., Yin, D., Tian, Y., Wei, X., & Zhou, M. (2025). Exploring Similarity between Neural and LLM Trajectories in Language Processing. arXiv preprint arXiv:2509.24307.

Xu, N., Zhang, Q., Du, C., Luo, Q., Qiu, X., Huang, X., & Zhang, M. (2025). Human-like conceptual representations emerge from language prediction. *Proceedings of the National Academy of Sciences*, 122(44), e2512514122. https://doi.org/10.1073/pnas.2512514122

Yao, Y., Yang, Y., Ma, X., Yang, D., Zhang, Z., Li, Z., & Zhao, H. (2025). How deep is love in LLMs' hearts? Exploring semantic size in human-like cognition. *arXiv preprint* arXiv:2503.00330. https://arxiv.org/abs/2503.00330

Zhang, Z. Y., Verma, A., Doshi-Velez, F., & Low, B. K. H. (2024). Understanding the relationship between prompts and response uncertainty in large language models. [Preprint]. *arXiv*. https://arxiv.org/abs/2407.14845. (LLMs internally gauge and respond to uncertainty in prompts, indicating genuine comprehension and probabilistic reasoning rather than simple pattern-matching.)